

## Continuous integration in gitlab

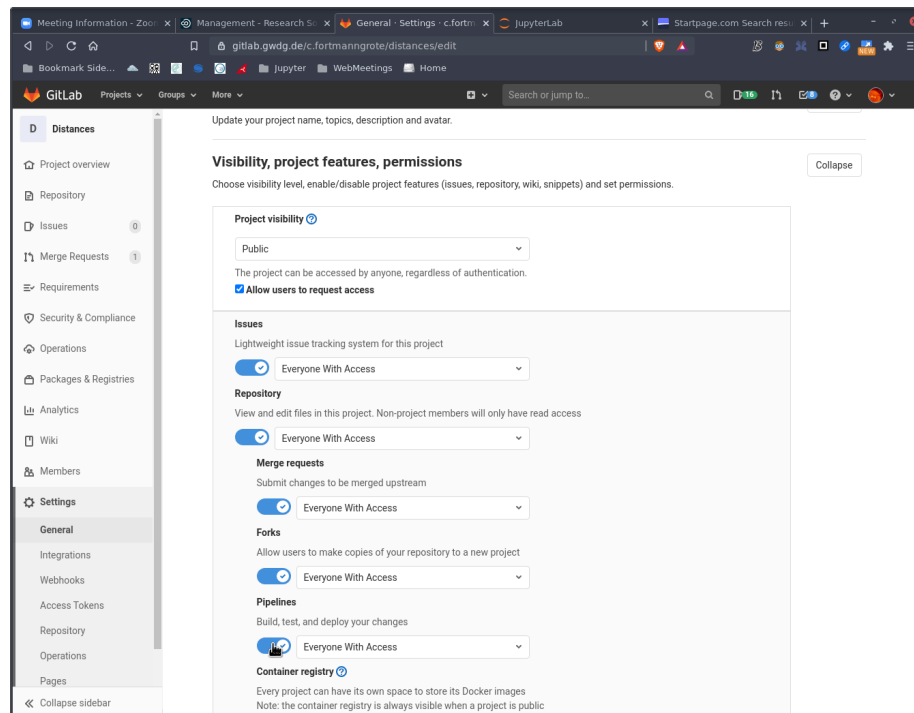
In this chapter, we will introduce the concept of "Continuous Integration", short "CI". CI means that every change to the codebase will be integrated in such a way that its effect on the test status, documentation, and released software packages become immediately visible. That last aspect, deployment of updated software is also referred to as "Continuous Deployment, CD".

Here, we will focus on automated testing and documentation. Therefore, we have to instruct the GitLab server to trigger a run of our test suite after every `git push`.

In the previous part, we have shown how to run your tests with `pytest`. Now we will go one step further and instruct gitlab to run the test *every time* a change is committed to the code base.

Detailed information about CI in gitlab can be found in the official gitlab documentation at <https://docs.gitlab.com/ee/ci/>. Here we only cover the basic steps to setup CI for your project.

In your web browser, navigate to the gitlab repository front page and open **Settings** from the left navigation panel. Expand the second item "Visibility, ..." and activate the toggle named "Pipelines".

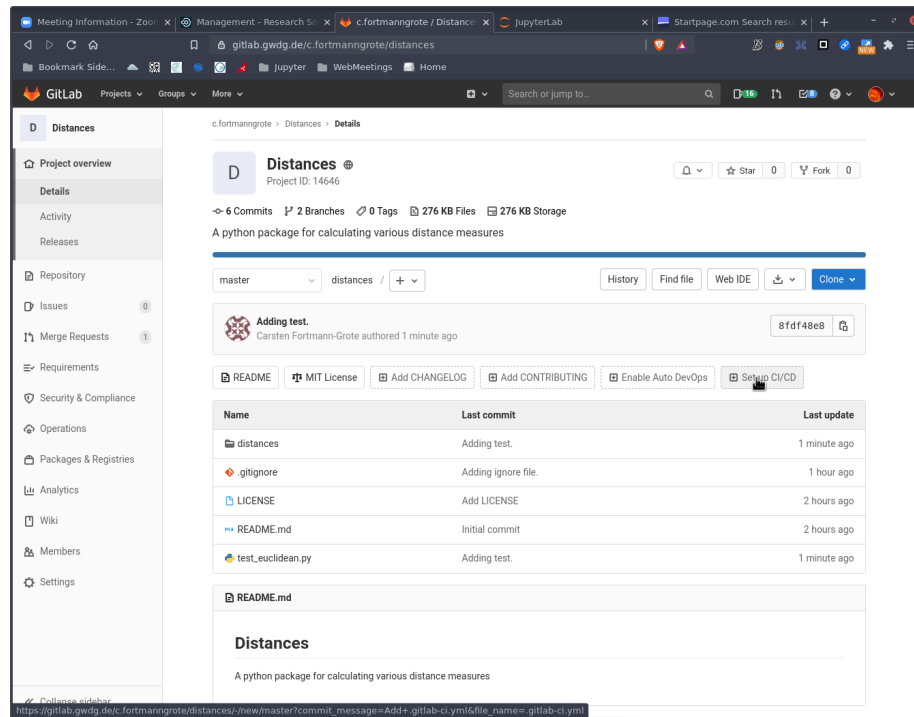


Don't forget to click **Save changes** at the bottom of the page.

## Activate CI through the `.gitlab-ci.yml` file.

Gitlab automatically scans the root directory for a file named `.gitlab-ci.yml`. If that file is detected, gitlab will parse this file and execute the instructions. The file must follow the `yaml` ("Yet Another Markup Language") syntax.

On the front page of your gitlab repo, click on **Setup CI/CD**



Paste the following lines into the text entry field:

```
# gitlab-ci yaml file for project "distances"
# The docker base image to use.
image: python:latest

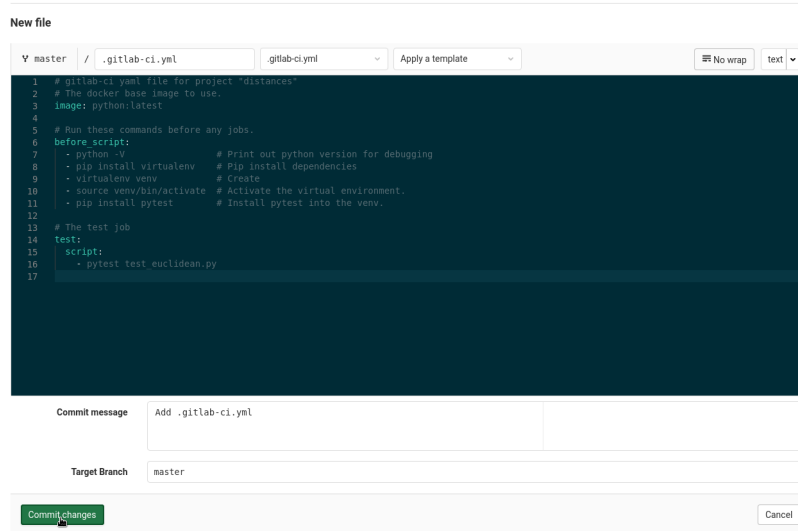
# Run these commands before any jobs.
before_script:
  - python -V # Print out python version for debugging
  - pip install virtualenv # Pip install dependencies
  - virtualenv venv # Create
  - source venv/bin/activate # Activate the virtual environment.
  - pip install pytest # Install pytest into the venv.

# The test job
test:
```

```
script:
  - pytest test_euclidean.py
```

This `yaml` file is structured as follows:

- The "base image": Gitlab-CI uses docker images to encapsulate the test environment.
- `before_script`: Contains commands to execute before any jobs are run.
- `test`: The `test` job runs the test code through `pytest`.
- More jobs can be defined.



New file

master / .gitlab-ci.yml gitlab-ci.yml Apply a template No wrap text

```
1 # gitlab-ci.yml file for project "distances"
2 # The docker base image to use.
3 image: python:latest
4
5 # Run these commands before any jobs.
6 before_script:
7   - python -V           # Print out python version for debugging
8   - pip install virtualenv # Pip install dependencies
9   - virtualenv venv      # Create
10  - source venv/bin/activate # Activate the virtual environment.
11  - pip install pytest    # Install pytest into the venv.
12
13 # The test job
14 test:
15   script:
16     - pytest test_euclidean.py
17
```

Commit message Add .gitlab-ci.yml

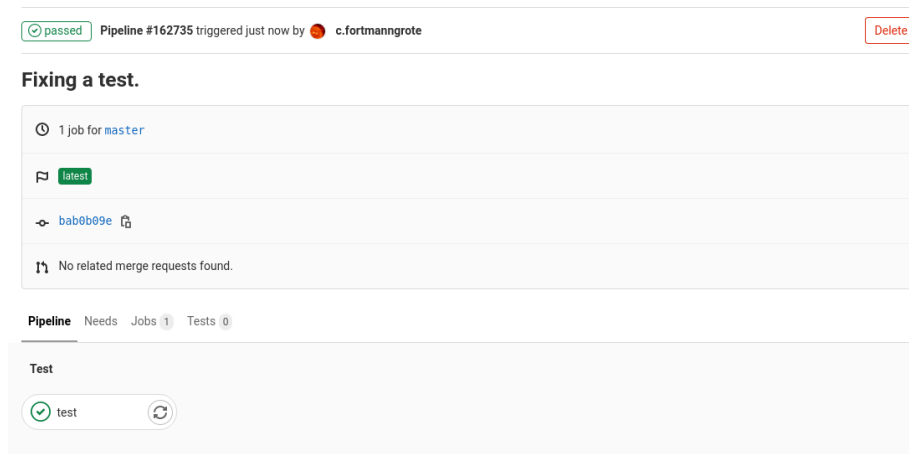
Target Branch master

Commit changes Cancel

Finalize the CI configuration by committing the changes to your repository. Provide a commit message and click the "Commit changes" button.

Since CI is activated (see above), this commit will trigger the first CI run. Check the process by navigating to `CI/CD`.

After the test is finished, you will see a test report page, similar to this one (it will look differently if the test failed...)



Now let's try out the automated test feature. We will implement an additional distance, the Manhattan distance. But this time, we'll do it the pro way: We will *first* add a test function for the new distance metric. Naturally, the test will fail. Subsequently, we will implement the new distance such that the test passes. This approach is called *test driven development* and is an established technique in software engineering. The value of writing the test first is that we have to get a very clear idea of the expected behaviour of our new function, even before we start implementing it. If we were given a list of requirements from our supervisor, we could add a test function for each requirement and then implement the code such that one test after the other succeeds. In the end, all requirements would be satisfied.

## Exercises

1. Make "atomic" commits, i.e. one change per commit.
2. Update your local repository's master branch (`git pull`).
3. Create and checkout a new branch named "manhattan".
4. Copy the file `test_euclidean.py` to a new file `test_manhattan.py`
5. Implement the test using the same pair of vectors as in `test_euclidean.py`.  
What is the expected Manhattan distance between vectors  $u$  and  $v$ ?
6. Run the new test and observe the test failure.
7. Implement the Manhattan distance in a new file `distances/manhattan.py`.
8. Add the new module to the distances package (Hint: Edit `distances/__init__.py`)
9. Add the new test to the CI configuration.

## Solutions:

1. Update `shell` `git checkout master` `git pull origin master`
2. New branch off master `shell` `git branch manhattan` `git`

```
checkout manhattan
```

3. Copy shell `cp test_euclidean.py test_manhattan.py` `git add test_manhattan.py` `git commit`

4. Implement the test in `test_manhattan.py`:

```
import distances

def test_manhattan():
    """ Test the manhattan distance calculation. """
    u = (2, -1)
    v = (-2, 2)

    assert distances.manhattan_distance(u, v) == 7.0
```

“shell git add git commit

5. Run the test shell `pytest test_manhattan.py`

As expected, the test will fail because the `distances` package has no definition of the `manhattan_distance()` function that is called in the `test_manhattan` test function.

6. Implement the new distance function.

```
# distances/manhattan.py
import math

def manhattan_distance(u, v):
    """
    Computes the Manhattan distance between two vectors `u` and `v`.

    The Euclidean distance between `u` and `v`, is defined as:

     $|u_1 - v_1| + \dots + |u_n - v_n|$ 

    Parameters
    -----
    u : list
        Input vector.
    v : list
        Input vector.

    Returns
    -----
    manhattan : double
        The Manhattan distance between vectors `u` and `v`.
    """
    distance = 0
```

```

for u_i, v_i in zip(u, v):
    distance += abs(u_i - v_i)**2

return distance

```

```

git add distances/manhattan.py
git commit

```

7. Add to package “python # distances/**init**.py  
from .euclidean import euclidean\_distance “
8. Edit .gitlab-ci.yml: “yml # gitlab-ci yml file for project "distances"  
# The docker base image to use. image: python:latest  
# Run these commands before any jobs. before\_script:
  - python -V # Print out python version for debugging
  - pip install virtualenv # Pip install dependencies
  - virtualenv venv # Create
  - source venv/bin/activate # Activate the virtual environment.
  - pip install pytest # Install pytest into the venv.
# The test job test: script: - pytest test\_euclidean.py - pytest  
test\_manhattan.py “

Add the two new files and one modified file to git:

```

$ git add test_manhattan.py
$ git add distances/manhattan.py
$ git add distances/__init__.py

```

And commit the changes:

```

$ git commit -m "Add manhattan distance and test."

```

**Q:** What is the function of the "-m" flag in the `git commit` command above?

## Register the new test module to CI

Finally, we have to register the new test to the CI system. Edit the file .gitlab-ci.yml as follows:

```

# gitlab-ci yml file for project "distances"
# The docker base image to use.
image: python:latest

# Run these commands before any jobs.
before_script:
- python -V # Print out python version for debugging
- pip install virtualenv # Pip install dependencies

```

```

- virtualenv venv          # Create
- source venv/bin/activate # Activate the virtual environment.
- pip install pytest       # Install pytest into the venv.

# The test job
test:
  script:
    - pytest test_euclidean.py
    - pytest test_manhattan.py

```

Add and commit this change:

```

$ git add .gitlab-ci.
$ git commit

```

and push all changes to the remote branch on gitlab:

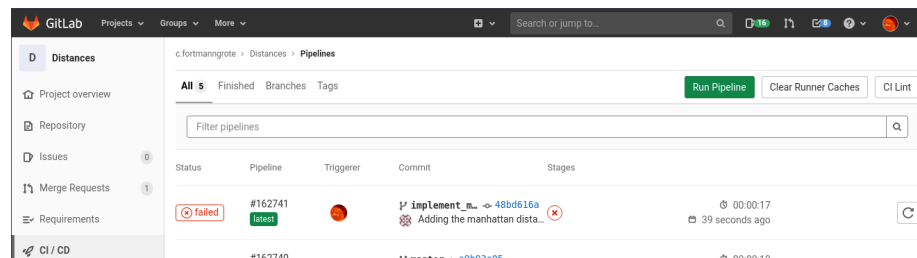
```

$ git push -u origin implement_manhattan_distance

```

## Check test status on gitlab

Now open your repository website and navigate to the CI summary page (click CI/CD in the navigation panel).



## What? The test failed!

Inspecting the test report, we find that the expected result and the result our code gave, do not agree:

```

76 ===== 1 passed, 1 warning in 0.01s =====
77 $ pytest test_manhattan.py
78 ===== test session starts =====
79 platform linux -- Python 3.9.0, pytest-6.1.2, py-1.9.0, pluggy-0.13.1
80 rootdir: /builds/c.fortmanngröte/distances
81 collected 1 item
82 test_manhattan.py F [100%]
83 ===== FAILURES =====
84 _____ test_manhattan _____
85     def test_manhattan():
86         """ Test the Manhattan distance calculation. """
87         u = (2, -1)
88         v = (-2, 2)
89
90     >     assert distances.manhattan_distance(u, v) == 7.0
91     E     assert 25 == 7.0
92     E         + where 25 = <function manhattan_distance at 0x7f7650b7a820>((2, -1), (-2,
93     E         +       where <function manhattan_distance at 0x7f7650b7a820> = distances.manhatt
94     E         an_distance
95     test_manhattan.py:8: AssertionError
96 ===== short test summary info =====
96 FAILED test_manhattan.py::test_manhattan - assert 25 == 7.0
97 ===== 1 failed in 0.02s =====

```

**Q:** Try to spot the error in `distances/manhattan.py` and fix it. Then recommit to the repository and observe the test run. Iterate this procedure until the test goes through.

## Building the documentation

The CI feature of gitlab not only allows us to run tests. In the `.gitlab-ci.yml`, we can define any number of "jobs" that will be run sequentially in a "pipeline". Typically, a complete CI pipeline consists of

- Installing dependencies
- Building the code
- Running the tests
- Building the documentation
- Prepare a software package for the users to download

Here, we will demonstrate how to build a reference manual for our code.

The reference manual will be built based on the information we put into the docstrings of our functions in our python modules. Take a look at `distances/euclidean.py`:

```
# distances/euclidean.py
```

```
import math
```



```
def euclidean_distance(u, v):
    """
    Computes the Euclidean distance between two vectors `u` and `v`.

    The Euclidean distance between `u` and `v`, is defined as:

    
$$\sqrt{(u_1 - v_1)^2 + \dots + (u_n - v_n)^2}$$


    Parameters
    -----
    u : list
        Input vector.
    v : list
        Input vector.

    Returns
    -----
    euclidean : double
        The Euclidean distance between vectors `u` and `v`.
    """
    distance = 0

    for u_i, v_i in zip(u, v):
        distance += (u_i - v_i) ** 2

    return math.sqrt(distance)
```

The part between a pair of triple-`"""` is the docstring. It describes what the function does and the type of input parameters as well as the return type.

This docstring can be parsed by special documentation builders and turned into a reference manual document. The output can be saved as a html document (online documentation) or as a pdf. Finally, we will see how this process can be run as a CI job and the documentation be served as a website on gitlab.

## Install and configure sphinx

The documentation parser is called `sphinx` and can be installed through `anaconda`. On the command line, run

```
$ conda install sphinx
```

or use your system's `anaconda` GUI.

As always, we will work in a separate branch:

```
$ git checkout -b doc
```

Create a new `doc` directory under the repository root directory,

```
$ mkdir doc
```

Your repository directory tree should look like this:

```
|--- .git/
|--- distances/
|    |--- euclidean.py
|    |--- manhattan.py
|    |--- __init__.py
|--- doc/
|--- test_euclidean.py
|--- test_manhattan.py
|--- .gitignore
|--- .gitlab-ci.yml
|--- LICENSE
|--- README.md
```

Change into the new directory

```
$ cd doc
```

To configure the documentation builder `sphinx`, use the `quick-conf` utility:

```
$ sphinx-quickstart
```

Answer the questions as follows:

```
> Separate source and build directories (y/n) [n]: y
> Project name: Distances
> Author name(s): <YOUR_NAME_HERE>
> Project release []:
> Project language [en]:
```

The `sphinx-quickstart` has created a number of new file in the current working directory:

```
$ ls -l
build/  make.bat  Makefile  source/
```

The files `Makefile` and `make.bat` contain all the detailed instruction for your systems `make` utility to actually build the documentation. We don't have to worry about these files. The directory `source/` will contain all the *static* source files from which to build the documentation. Finally, the build documentation will be saved under the `build/` directory.

Next, `cd` into the `source/` directory:

```
$ cd source
```

And open the file `conf.py` in your editor. Locate the section heading `#-- Path setup ---...` and change the code below:

```
import os
import sys
sys.path.insert(0, os.path.abspath('../..'))
from distances import *
```

This adds our `distances` package to the documentation system.

Further down, under `#-- General configuration -- ...`, add the `autodoc` extension:

```
extensions = [
    "sphinx.ext.autodoc",
    "sphinx.ext.napoleon",
]
```

The `autodoc` extension is responsible for parsing the docstrings and formatting them into a reference manual. The `napoleon` extension adds support for Google and numpy documentation syntax used in our code (see [this link](#)) for more info.

Save and close the file.

## Manually building the documentation

At this point, we can already build the skeleton of our documentation, i.e. the layout and general markup. There is no content yet...

Change back to the parent directory

```
$ cd ..
```

and run the command

```
$ make html
```

This will build the documentation website html document and save it under `build/html`. You can open the page by running

```
$ python -m http.server --directory build/html
```

Then open the website `http://localhost:8000` in your browser.

You should see a static website similar to this:

## Adding the documentation content

Now it's time to add some content to our documentation and link in the docstrings from the code.

The file `doc/source/index.rst` is the root document of the documentation. It is formatted in the `ReStructuredText` markup language. E.g. top level headings are underlined by `===`, directives start with `..` and so on. See the `ReStructuredText` manual for details. To add our modules' documentation, we first have to convert their docstrings into `ReStructuredText`: From the `doc/` directory, run

```
$ sphinx-apidoc -f -o source ../distances
```

This will create two new files in the `source/` directory, `modules.rst` and `distances.rst`. These contain the formatted version of our docstrings. Now, we can easily include the distances documentation in the `index.rst` by adding one line after the `.. toctree` directive. Your `index.rst` should look like this:

```
.. doc/source/index.rst
```

```
Welcome to Distances's documentation!
```

```
=====
```

```
.. toctree::
   :maxdepth: 2
   :caption: Contents:
```

```
distances
```

```
Indices and tables
```

```
=====
```

```
* :ref:`genindex`
```

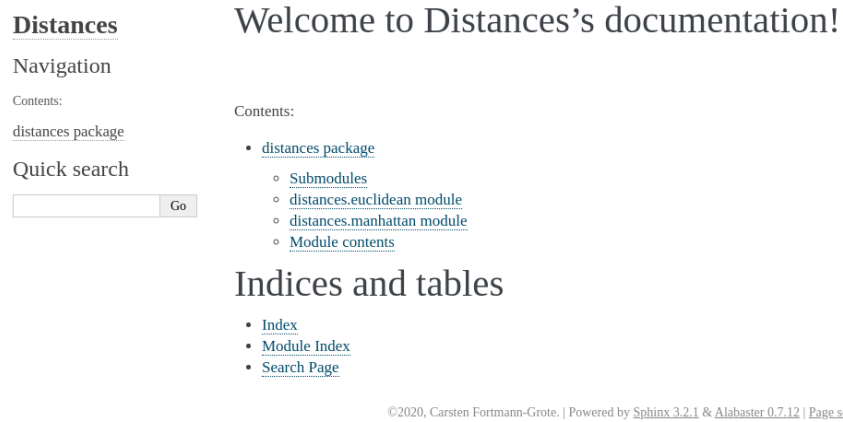
```
* :ref:`modindex`
* :ref:`search`
```

Build the documentation again:

```
$ make html
```

Reload the documentation website at <http://localhost:8000> (supposing that the webserver is still running, otherwise restart it, see above).

You should now see a beautiful documentation website similar to this one:



Take some time to browse around your documentation site and try to mentally connect the contents of the `.rst` files and the elements and links in the webpage.

## Let gitlab-CI build the documentation: gitlab-pages

So far we have *manually* built the documentation. What we want is that each time the code and/or the documentation changes, the manual is rebuilt to reflect these changes. We need to define a new CI job in our pipeline.

Change back to the repository root directory and edit the file `.gitlab-ci.yml`.

```
# gitlab-ci yml file for project "distances"
# The docker base image to use.
image: python:latest

# Run these commands before any jobs.
before_script:
  - python -V # Print out python version for debugging
  - pip install virtualenv # Pip install dependencies
  - virtualenv venv # Create
```

```

- source venv/bin/activate # Activate the virtual environment.
- pip install pytest      # Install pytest into the venv.
- pip install sphinx      # Install sphinx into the venv.

# The test job
test:
  script:
    - pytest test_euclidean.py
    - pytest test_manhattan.py

pages:
  script:
    - cd doc
    - sphinx-apidoc -f -o source ../distances
    - make html
    - cp -av build/html ../public
  artifacts:
    paths:
      - public

```

We add the **pages** job at the end. It contains all steps to take to build the documentation. The last step is to copy the built documentation website to a directory **pages** under the repository root. The **artifacts** directive ensures that the **public** directory will be maintained when the site is exposed through a webserver.

## Push to gitlab

We have now configured our repository to automatically build the reference manual. Let's add all needed new files to git and push to gitlab.

```

$ git add doc/Makefile
$ git add doc/make.bat
$ git add doc/source/conf.py
$ git add doc/source/index.rst
$ git add .gitlab-ci.yml

$ git commit
$ git push -u origin doc

```

The last command pushes our changes to a new remote branch **doc**.

Observe the status of your CI pipeline under CI/CD -> Pipelines. Note that the pipeline now contains two Jobs, one for testing and one for building the documentation. After the second job succeeds, you should be able to open your online documentation site at <http://.pages.gwdg.de/distances>.

## Build documentation only for the master branch

To make sure that the documentation is built only on the master branch, we add one more directive to `.gitlab-ci.yml`:

```
# gitlab-ci yaml file for project "distances"
# The docker base image to use.
image: python:latest

# Run these commands before any jobs.
before_script:
  - python -V           # Print out python version for debugging
  - pip install virtualenv # Pip install dependencies
  - virtualenv venv      # Create
  - source venv/bin/activate # Activate the virtual environment.
  - pip install pytest   # Install pytest into the venv.
  - pip install sphinx    # Install pytest into the venv.

# The test job
test:
  script:
    - pytest test_euclidean.py
    - pytest test_manhattan.py

pages:
  script:
    - cd doc
    - sphinx-apidoc -f -o source ../distances
    - make html
    - cp -av build/html ../public
artifacts:
  paths:
    - public
only:
  - master
```

Add, commit, and push again and then create a merge request to merge the `doc` branch into `master`. Note that the CI pipeline is also run on the Merge request, so you can detect possible issues already before the final merge.

## Add test and documentation badges to README

As a final gimmick, let's be hip and show off our test status on the repository's front page. Follow these steps:

- Navigate to your project's Settings > General > Badges.
- Under Name, enter Pipeline Status.

- Under Link, enter the following URL: `https://gitlab.com/%7Bproject_path%7D/-/commits/%7Bdefault_branch%7D`
- Under Badge image URL, enter the following URL: `https://gitlab.com/%7Bproject_path%7D/badges/%7Bdefault_branch%7D/pipeline.svg`
- Submit the badge by clicking the Add badge button.

The pipeline status badge will then appear on top of the repository's front page: